

Soft Actor-Critic (SAC)

이동민

삼성전자 서울대 공동연구소
Jul 19, 2019

Outline

- Soft Actor-Critic (SAC)
 - Learning process
 - Hyperparameter
 - Main loop
 - Train model
 - Train & TensorboardX
 - Learning curve & Test
- Comparison of algorithms for continuous action

Soft Actor-Critic (SAC)

- SAC Algorithm

Algorithm 1 Soft Actor-Critic

Input: θ_1, θ_2, ϕ

$\bar{\theta}_1 \leftarrow \theta_1, \bar{\theta}_2 \leftarrow \theta_2$

$\mathcal{D} \leftarrow \emptyset$

for each iteration **do**

for each environment step **do**

$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$

$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$

$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$

end for

for each gradient step **do**

$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i)$ for $i \in \{1, 2\}$

$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$

$\alpha \leftarrow \alpha - \lambda \hat{\nabla}_\alpha J(\alpha)$

$\bar{\theta}_i \leftarrow \tau \theta_i + (1 - \tau) \bar{\theta}_i$ for $i \in \{1, 2\}$

end for

end for

Output: θ_1, θ_2, ϕ

▷ Initial parameters

▷ Initialize target network weights

▷ Initialize an empty replay pool

▷ Sample action from the policy

▷ Sample transition from the environment

▷ Store the transition in the replay pool

▷ Update the Q-function parameters

▷ Update policy weights

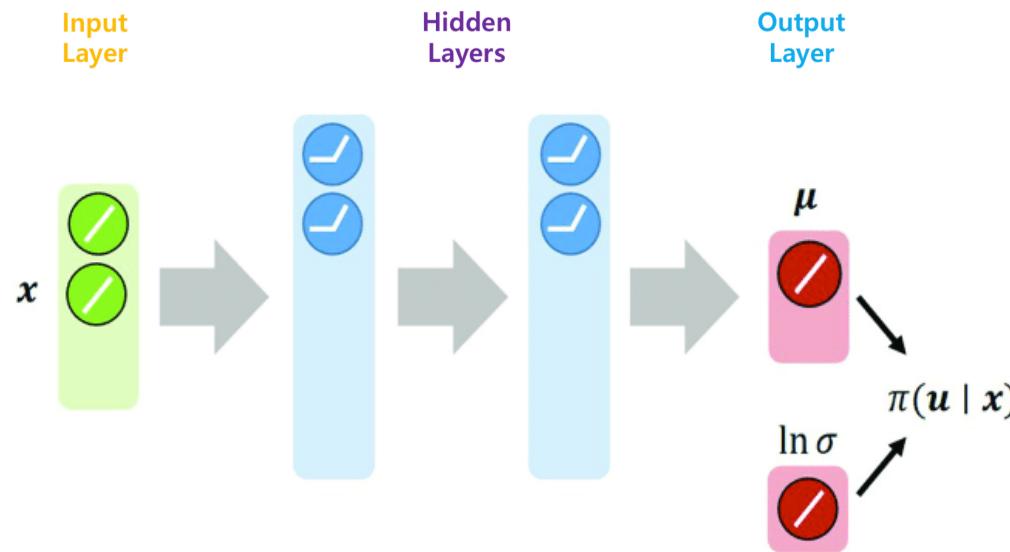
▷ Adjust temperature

▷ Update target network weights

▷ Optimized parameters

Soft Actor-Critic (SAC)

- Actor network



source : [Link](#)

Soft Actor-Critic (SAC)

- Actor network

```
class Actor(nn.Module):
    def __init__(self, state_size, action_size, args, log_std_min=-20, log_std_max=2):
        super(Actor, self).__init__()
        self.log_std_min = log_std_min
        self.log_std_max = log_std_max

        self.fc1 = nn.Linear(state_size, args.hidden_size)
        self.fc2 = nn.Linear(args.hidden_size, args.hidden_size)

        self.fc3 = nn.Linear(args.hidden_size, action_size)
        self.fc4 = nn.Linear(args.hidden_size, action_size)

    def forward(self, x):
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))

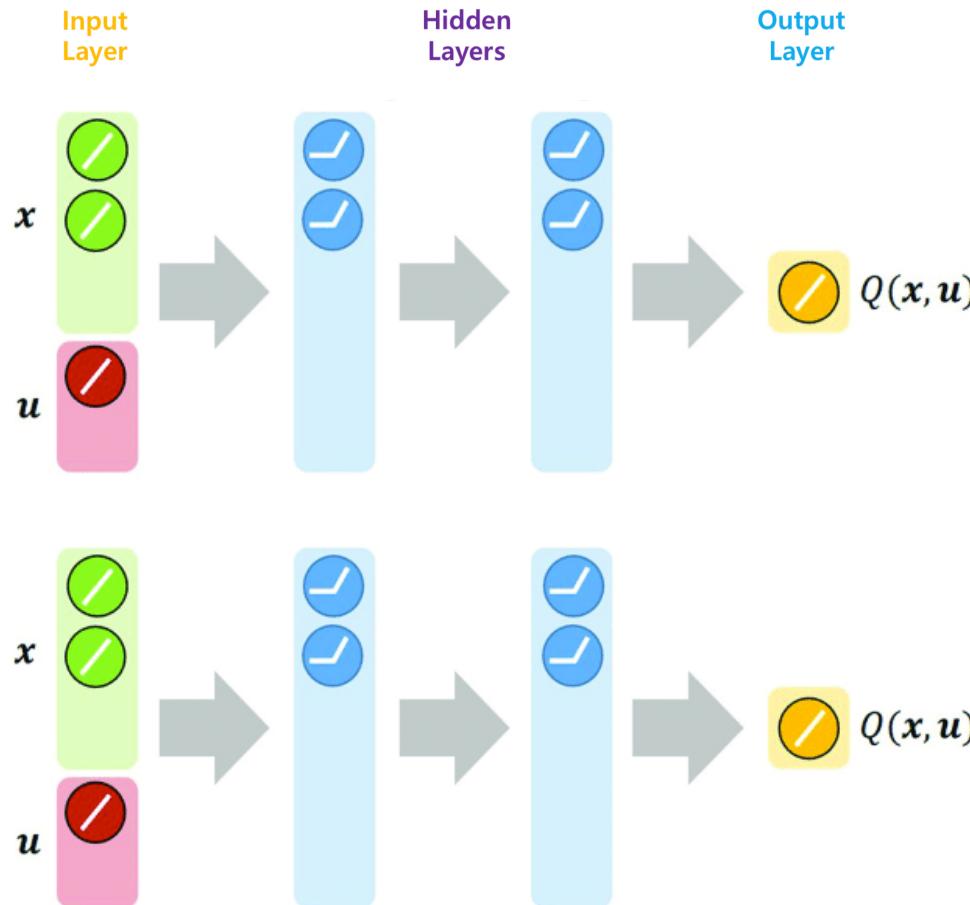
        mu = self.fc3(x)
        log_std = self.fc4(x)

        log_std = torch.clamp(log_std, min=self.log_std_min, max=self.log_std_max)
        std = torch.exp(log_std)

        return mu, std
```

Soft Actor-Critic (SAC)

- Critic network



source : [Link](#)

Soft Actor-Critic (SAC)

- Critic network

```
class Critic(nn.Module):
    def __init__(self, state_size, action_size, args):
        super(Critic, self).__init__()

        # Q1 architecture
        self.fc1 = nn.Linear(state_size + action_size, args.hidden_size)
        self.fc2 = nn.Linear(args.hidden_size, args.hidden_size)
        self.fc3 = nn.Linear(args.hidden_size, 1)

        # Q2 architecture
        self.fc4 = nn.Linear(state_size + action_size, args.hidden_size)
        self.fc5 = nn.Linear(args.hidden_size, args.hidden_size)
        self.fc6 = nn.Linear(args.hidden_size, 1)

    def forward(self, states, actions):
        x = torch.cat([states, actions], dim=1)

        x1 = torch.relu(self.fc1(x))
        x1 = torch.relu(self.fc2(x1))
        q_value1 = self.fc3(x1)

        x2 = torch.relu(self.fc4(x))
        x2 = torch.relu(self.fc5(x2))
        q_value2 = self.fc6(x2)

        return q_value1, q_value2
```



Soft Actor-Critic (SAC)

- Learning process
 1. 상태에 따른 행동 선택
 2. 환경에서 선택한 행동으로 한 time step을 진행한 후, 다음 상태와 보상을 받음
 3. Sample (s, a, r, s') 을 replay buffer에 저장
 4. Replay buffer에서 랜덤으로 sample을 추출
 5. 추출한 sample로 Actor & Critic network 업데이트
 6. Temperature parameter alpha 업데이트
 7. Critic에 대해 soft target 업데이트

Learning process

1. 상태에 따른 행동 선택

```
140     mu, std = actor(torch.Tensor(state))  
141     action = get_action(mu, std)
```

```
def get_action(mu, std):  
    normal = Normal(mu, std)  
    z = normal.rsample() # reparameterization trick (mean + std * N(0,1))  
    action = torch.tanh(z)  
  
    return action.data.numpy()
```

- `normal.rsample()`

```
[docs] def rsample(self, sample_shape=torch.Size()):  
    shape = self._extended_shape(sample_shape)  
    eps = _standard_normal(shape, dtype=self.loc.dtype, device=self.loc.device)  
    return self.loc + eps * self.scale
```

Learning process

2. 환경에서 선택한 행동으로 한 time step을 진행한 후, 다음 상태와 보상을 받음

```
143 | | | next_state, reward, done, _ = env.step(action)
```

3. Sample (s, a, r, s') 을 replay buffer에 저장

```
123 | | | replay_buffer = deque(maxlen=10000)  
  
146 | | | mask = 0 if done else 1  
147 | | |  
148 | | | replay_buffer.append((state, action, reward, next_state, mask))
```

4. Replay buffer에서 랜덤으로 sample을 추출 (Batch size : 64)

```
154 | | | mini_batch = random.sample(replay_buffer, args.batch_size)
```

Learning process

5. 추출한 sample로 Actor & Critic network 업데이트

- Critic Loss

$$J_Q(\theta) = (\underbrace{Q_\theta(s, a)}_{\text{Prediction}} - \underbrace{(r + \gamma(Q_{\theta^-}(s', a') - \alpha \log \pi_\phi(a'|s')))}_{\text{Target}})^2$$

```
48     # update critic
49     criterion = torch.nn.MSELoss()
50
51     # get Q-values using two Q-functions to mitigate overestimation bias
52     q_value1, q_value2 = critic(torch.Tensor(states), actions)
53
54     # get target
55     mu, std = actor(torch.Tensor(next_states))
56     next_policy, next_log_policy = eval_action(mu, std)
57     target_next_q_value1, target_next_q_value2 = target_critic(torch.Tensor(next_states), next_policy)
58
59     min_target_next_q_value = torch.min(target_next_q_value1, target_next_q_value2)
60     min_target_next_q_value = min_target_next_q_value.squeeze(1) - alpha * next_log_policy.squeeze(1)
61     target = rewards + masks * args.gamma * min_target_next_q_value
62
63     critic_loss1 = criterion(q_value1.squeeze(1), target.detach()) # Equation 5
64     critic_optimizer.zero_grad()
65     critic_loss1.backward()
66     critic_optimizer.step()
67
68     critic_loss2 = criterion(q_value2.squeeze(1), target.detach()) # Equation 5
69     critic_optimizer.zero_grad()
70     critic_loss2.backward()
71     critic_optimizer.step()
```

Learning process

5. 추출한 sample로 Actor & Critic network 업데이트

- Actor Loss

$$J_{\pi}(\phi) = \frac{1}{N} \sum \alpha \log \pi_{\phi}(f_{\phi}(\epsilon; s) | s) - Q_{\theta}(s, f_{\phi}(\epsilon; s))$$

```
73     # update actor
74     mu, std = actor(torch.Tensor(states))
75     policy, log_policy = eval_action(mu, std)
76
77     q_value1, q_value2 = critic(torch.Tensor(states), policy)
78     min_q_value = torch.min(q_value1, q_value2)
79
80     actor_loss = ((alpha * log_policy) - min_q_value).mean() # Equation 9
81     actor_optimizer.zero_grad()
82     actor_loss.backward()
83     actor_optimizer.step()
```

Learning process

6. Temperature parameter alpha 업데이트

- Alpha Loss

$$J(\alpha) = -\frac{1}{N} \sum \alpha (\log \pi_{\phi}(f_{\phi}(\epsilon; s) | s) + \bar{\mathcal{H}})$$

```
85     # update alpha
86     alpha_loss = -(log_alpha * (log_policy + target_entropy).detach()).mean() # Equation 18
87     alpha_optimizer.zero_grad()
88     alpha_loss.backward()
89     alpha_optimizer.step()
90
91     alpha = torch.exp(log_alpha)
92
93     return alpha
```

- log_alpha : alpha를 (0,1] 사이의 값으로 만들어주기 위해 사용

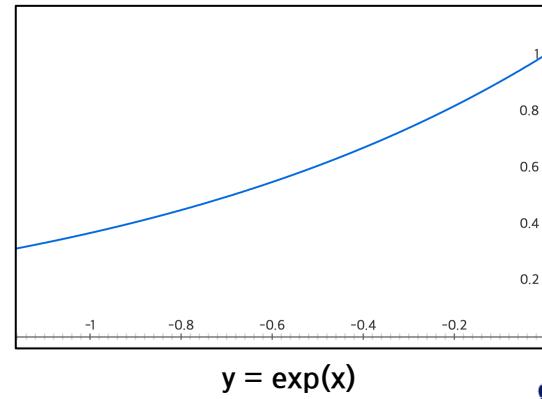
```
log_alpha tensor([0.], requires_grad=True)
alpha tensor([1.], grad_fn=<ExpBackward>)

log_alpha tensor([-0.0001], requires_grad=True)
alpha tensor([0.9999], grad_fn=<ExpBackward>)

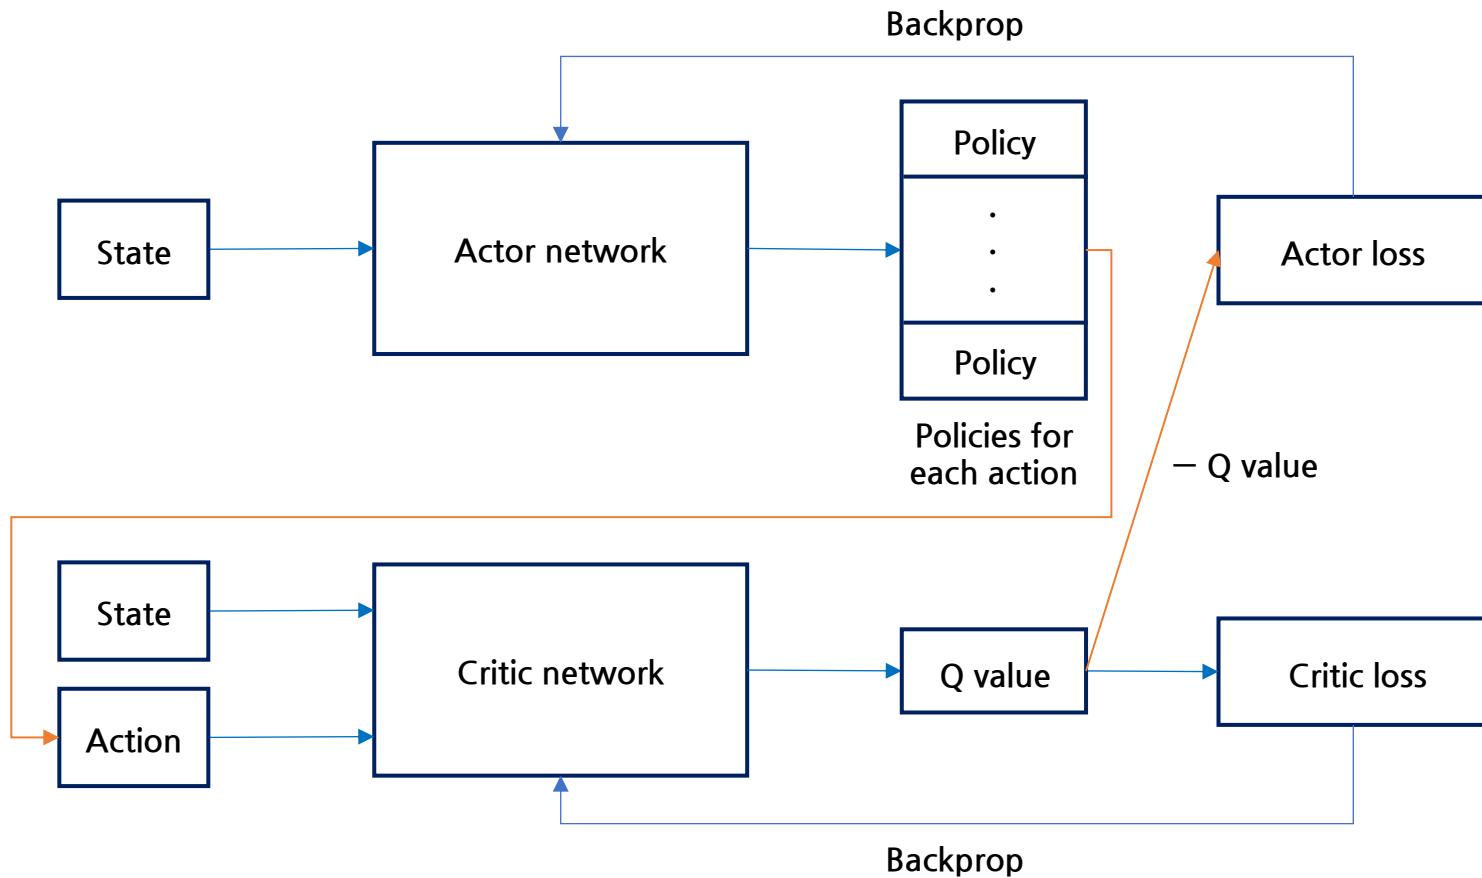
log_alpha tensor([-0.0002], requires_grad=True)
alpha tensor([0.9998], grad_fn=<ExpBackward>)

log_alpha tensor([-0.0003], requires_grad=True)
alpha tensor([0.9997], grad_fn=<ExpBackward>)

log_alpha tensor([-0.0004], requires_grad=True)
alpha tensor([0.9996], grad_fn=<ExpBackward>)
```



Learning process



Learning process

7. Critic에 대해 soft target 업데이트

- Initialize target model

```
113 |     hard_target_update(critic, target_critic)
```

```
def hard_target_update(net, target_net):  
    target_net.load_state_dict(net.state_dict())
```

- Soft target update ($\tau : 0.005$)

$$\theta^{Q'} \leftarrow \tau\theta^Q + (1 - \tau)\theta^{Q'}$$

$$\phi^{\pi'} \leftarrow \tau\phi^\pi + (1 - \tau)\phi^{\pi'}$$

```
161 |     |     |     |     soft_target_update(critic, target_critic, args.tau)
```

```
def soft_target_update(net, target_net, tau):  
    for param, target_param in zip(net.parameters(), target_net.parameters()):  
        target_param.data.copy_(tau * param.data + (1.0 - tau) * target_param.data)
```

Hyperparameter

```
15 parser = argparse.ArgumentParser()
16 parser.add_argument('--env_name', type=str, default="Pendulum-v0")
17 parser.add_argument('--load_model', type=str, default=None)
18 parser.add_argument('--save_path', default='./save_model/', help='')
19 parser.add_argument('--render', action="store_true", default=False)
20 parser.add_argument('--gamma', type=float, default=0.99)
21 parser.add_argument('--hidden_size', type=int, default=64)
22 parser.add_argument('--batch_size', type=int, default=64)
23 parser.add_argument('--actor_lr', type=float, default=1e-3)
24 parser.add_argument('--critic_lr', type=float, default=1e-3)
25 parser.add_argument('--alpha_lr', type=float, default=1e-4)
26 parser.add_argument('--tau', type=float, default=0.005)
27 parser.add_argument('--max_iter_num', type=int, default=1000)
28 parser.add_argument('--log_interval', type=int, default=10)
29 parser.add_argument('--goal_score', type=int, default=-300)
30 parser.add_argument('--logdir', type=str, default='./logs',
31 | | | | help='tensorboardx logs directory')
32 args = parser.parse_args()
```

Main loop

- Initialization

- Seed - random number 고정
- Actor & Critic network
- Target critic network
- Actor & Critic optimizer
- Hard target update
- Automatic entropy tuning
 - Target entropy
 - Log alpha, alpha
 - Alpha optimizer
- TensorboardX
- Replay buffer
- Recent rewards

```
96     def main():
97         env = gym.make(args.env_name)
98         env.seed(500)
99         torch.manual_seed(500)
100
101        state_size = env.observation_space.shape[0]
102        action_size = env.action_space.shape[0]
103        print('state size:', state_size)
104        print('action size:', action_size)
105
106        actor = Actor(state_size, action_size, args)
107        critic = Critic(state_size, action_size, args)
108        target_critic = Critic(state_size, action_size, args)
109
110        actor_optimizer = optim.Adam(actor.parameters(), lr=args.actor_lr)
111        critic_optimizer = optim.Adam(critic.parameters(), lr=args.critic_lr)
112
113        hard_target_update(critic, target_critic)
114
115        # initialize automatic entropy tuning
116        target_entropy = -torch.prod(torch.Tensor(action_size)).item()
117        log_alpha = torch.zeros(1, requires_grad=True)
118        alpha = torch.exp(log_alpha)
119        alpha_optimizer = optim.Adam([log_alpha], lr=args.alpha_lr)
120
121        writer = SummaryWriter(args.logdir)
122
123        replay_buffer = deque(maxlen=10000)
124        recent_rewards = deque(maxlen=100)
125        steps = 0
```

Main loop

- Episode 진행

- 상태에 따른 행동 선택
- 다음 상태와 보상을 받음
- Replay buffer에 저장

```
127     for episode in range(args.max_iter_num):
128         done = False
129         score = 0
130
131         state = env.reset()
132         state = np.reshape(state, [1, state_size])
133
134         while not done:
135             if args.render:
136                 env.render()
137
138             steps += 1
139
140             mu, std = actor(torch.Tensor(state))
141             action = get_action(mu, std)
142
143             next_state, reward, done, _ = env.step(action)
144
145             next_state = np.reshape(next_state, [1, state_size])
146             mask = 0 if done else 1
147
148             replay_buffer.append((state, action, reward, next_state, mask))
149
150             state = next_state
151             score += reward
```

Main loop

- Episode 진행
 - Replay buffer에서 랜덤으로 64개의 sample을 추출 → Mini batch
 - Train model
 - Critic에 대해 Soft target update

```
153     if steps > args.batch_size:  
154         mini_batch = random.sample(replay_buffer, args.batch_size)  
155  
156         actor.train(), critic.train(), target_critic.train()  
157         alpha = train_model(actor, critic, target_critic, mini_batch,  
158                             actor_optimizer, critic_optimizer, alpha_optimizer,  
159                             target_entropy, log_alpha, alpha)  
160  
161         soft_target_update(critic, target_critic, args.tau)  
162  
163     if done:  
164         recent_rewards.append(score)
```

Main loop

- Print & Visualize log
- Termination : 최근 100개의 episode의 평균 score가 -300보다 크다면
 - Save model
 - 학습 종료

```
166     if episode % args.log_interval == 0:  
167         print('{}) episode | score_avg: {:.2f}'.format(episode, np.mean(recent_rewards)))  
168         writer.add_scalar('log/score', float(score), episode)  
169  
170     if np.mean(recent_rewards) > args.goal_score:  
171         if not os.path.isdir(args.save_path):  
172             os.makedirs(args.save_path)  
173  
174         ckpt_path = args.save_path + 'model.pth'  
175         torch.save(actor.state_dict(), ckpt_path)  
176         print('Recent rewards exceed -300. So end')  
177         break
```

Train model

- Mini batch → Numpy array
- Mini batch에 있는 64개의 sample들을 각각 나눔
 - state - (64, 3)
 - action - (64, 1)
 - reward - (64)
 - next_state - (64, 3)
 - mask - (64)

```
34     def train_model(actor, critic, target_critic, mini_batch,
35     |         |         actor_optimizer, critic_optimizer, alpha_optimizer,
36     |         |         target_entropy, log_alpha, alpha):
37     |         mini_batch = np.array(mini_batch)
38     |         states = np.vstack(mini_batch[:, 0])
39     |         actions = list(mini_batch[:, 1])
40     |         rewards = list(mini_batch[:, 2])
41     |         next_states = np.vstack(mini_batch[:, 3])
42     |         masks = list(mini_batch[:, 4])
43
44     |         actions = torch.Tensor(actions).squeeze(1)
45     |         rewards = torch.Tensor(rewards).squeeze(1)
46     |         masks = torch.Tensor(masks)
```



Train model

- Prediction
- Target

```
48     # update critic
49     criterion = torch.nn.MSELoss()
50
51     # get Q-values using two Q-functions to mitigate overestimation bias
52     q_value1, q_value2 = critic(torch.Tensor(states), actions)
53
54     # get target
55     mu, std = actor(torch.Tensor(next_states))
56     next_policy, next_log_policy = eval_action(mu, std)
57     target_next_q_value1, target_next_q_value2 = target_critic(torch.Tensor(next_states), next_policy)
58
59     min_target_next_q_value = torch.min(target_next_q_value1, target_next_q_value2)
60     min_target_next_q_value = min_target_next_q_value.squeeze(1) - alpha * next_log_policy.squeeze(1)
61     target = rewards + masks * args.gamma * min_target_next_q_value
```

Train model

- Prediction
 - `q_value1, 2 - (64, 1)`
- Target

```
48      # update critic
49      criterion = torch.nn.MSELoss()
50
51      # get Q-values using two Q-functions to mitigate overestimation bias
52      q_value1, q_value2 = critic(torch.Tensor(states), actions)
```

Train model

- Prediction
 - **q_value1, 2 - (64, 1)**
- Target
 - **mu, std - (64, 1)**

```
54     # get target
55     mu, std = actor(torch.Tensor(next_states))
56     next_policy, next_log_policy = eval_action(mu, std) ←————
57     target_next_q_value1, target_next_q_value2 = target_critic(torch.Tensor(next_states), next_policy)
58
59     min_target_next_q_value = torch.min(target_next_q_value1, target_next_q_value2)
60     min_target_next_q_value = min_target_next_q_value.squeeze(1) - alpha * next_log_policy.squeeze(1)
61     target = rewards + masks * args.gamma * min_target_next_q_value
```

Train model

```
def eval_action(mu, std, epsilon=1e-6):
    normal = Normal(mu, std)
    z = normal.rsample() # reparameterization trick (mean + std * N(0,1))
    action = torch.tanh(z)
    log_prob = normal.log_prob(z)

    # Enforcing Action Bounds
    log_prob -= torch.log(1 - action.pow(2) + epsilon)
    log_policy = log_prob.sum(1, keepdim=True)

    return action, log_policy
```

$$\log \pi_{\phi}(a|s) = -\frac{(a - \mu_{\phi}(s))^2}{2\sigma^2} - \log \sigma - \log \sqrt{2\pi}$$

```
[docs] def log_prob(self, value):
    if self._validate_args:
        self._validate_sample(value)
    # compute the variance
    var = (self.scale ** 2)
    log_scale = math.log(self.scale) if isinstance(self.scale, Number) else
    self.scale.log()
    return -((value - self.loc) ** 2) / (2 * var) - log_scale - math.log(math.sqrt(2 *
    math.pi))
```

Train model

- Prediction
 - `q_value1, 2 - (64, 1)`
- Target
 - `mu, std - (64, 1)`

C Enforcing Action Bounds

We use an unbounded Gaussian as the action distribution. However, in practice, the actions needs to be bounded to a finite interval. To that end, we apply an invertible squashing function (\tanh) to the Gaussian samples, and employ the change of variables formula to compute the likelihoods of the bounded actions. In the other words, let $\mathbf{u} \in \mathbb{R}^D$ be a random variable and $\mu(\mathbf{u}|\mathbf{s})$ the corresponding density with infinite support. Then $\mathbf{a} = \tanh(\mathbf{u})$, where \tanh is applied elementwise, is a random variable with support in $(-1, 1)$ with a density given by

$$\pi(\mathbf{a}|\mathbf{s}) = \mu(\mathbf{u}|\mathbf{s}) \left| \det \left(\frac{d\mathbf{a}}{d\mathbf{u}} \right) \right|^{-1}. \quad (25)$$

Since the Jacobian $d\mathbf{a}/d\mathbf{u} = \text{diag}(1 - \tanh^2(\mathbf{u}))$ is diagonal, the log-likelihood has a simple form

$$\log \pi(\mathbf{a}|\mathbf{s}) = \log \mu(\mathbf{u}|\mathbf{s}) - \sum_{i=1}^D \log (1 - \tanh^2(u_i)), \quad (26)$$

where u_i is the i^{th} element of \mathbf{u} .

Train model

- Prediction
 - `q_value1, 2 - (64, 1)`
- Target
 - `mu, std - (64, 1)`

```
def eval_action(mu, std, epsilon=1e-6):
    normal = Normal(mu, std)
    z = normal.rsample() # reparameterization trick (mean + std * N(0,1))
    action = torch.tanh(z)
    log_prob = normal.log_prob(z)

    # Enforcing Action Bounds
    log_prob -= torch.log(1 - action.pow(2) + epsilon)
    log_policy = log_prob.sum(1, keepdim=True)

    return action, log_policy
```

$$\log \pi(\mathbf{a}|\mathbf{s}) = \log \mu(\mathbf{u}|\mathbf{s}) - \sum_{i=1}^D \log (1 - \tanh^2(u_i))$$

Train model

- Prediction
 - `q_value1, 2 - (64, 1)`
- Target
 - `mu, std - (64, 1)`
 - `next_policy, next_log_policy - (64, 1)`
 - `target_next_q_value1, 2 - (64, 1)`
 - `min_target_next_q_value - (64)`
 - `target - (64)`

```
54     # get target
55     mu, std = actor(torch.Tensor(next_states))
56     next_policy, next_log_policy = eval_action(mu, std)
57     target_next_q_value1, target_next_q_value2 = target_critic(torch.Tensor(next_states), next_policy)
58
59     min_target_next_q_value = torch.min(target_next_q_value1, target_next_q_value2)
60     min_target_next_q_value = min_target_next_q_value.squeeze(1) - alpha * next_log_policy.squeeze(1)
61     target = rewards + masks * args.gamma * min_target_next_q_value
```

Train model

- Update critic - MSE Loss

$$\circ \quad J_Q(\theta) = (\underbrace{Q_\theta(s, a)}_{\text{Prediction}} - \underbrace{(r + \gamma(Q_{\theta^-}(s', a') - \alpha \log \pi_\phi(a'|s')))}_{\text{Target}})^2$$

```
48     # update critic
49     criterion = torch.nn.MSELoss()
50
51     # get Q-values using two Q-functions to mitigate overestimation bias
52     q_value1, q_value2 = critic(torch.Tensor(states), actions)
53
54     # get target
55     mu, std = actor(torch.Tensor(next_states))
56     next_policy, next_log_policy = eval_action(mu, std)
57     target_next_q_value1, target_next_q_value2 = target_critic(torch.Tensor(next_states), next_policy)
58
59     min_target_next_q_value = torch.min(target_next_q_value1, target_next_q_value2)
60     min_target_next_q_value = min_target_next_q_value.squeeze(1) - alpha * next_log_policy.squeeze(1)
61     target = rewards + masks * args.gamma * min_target_next_q_value
62
63     critic_loss1 = criterion(q_value1.squeeze(1), target.detach()) # Equation 5
64     critic_optimizer.zero_grad()
65     critic_loss1.backward()
66     critic_optimizer.step()
67
68     critic_loss2 = criterion(q_value2.squeeze(1), target.detach()) # Equation 5
69     critic_optimizer.zero_grad()
70     critic_loss2.backward()
71     critic_optimizer.step()
```

Train model

- Update actor
 - **mu, std** - (64, 1)
 - **policy, log_policy** - (64, 1)
 - **q_value1, 2** - (64, 1)
 - **min_q_value** - (64, 1)
 - **alpha** - (1)
 - $J_{\pi}(\phi) = \frac{1}{N} \sum \alpha \log \pi_{\phi}(f_{\phi}(\epsilon; s) | s) - Q_{\theta}(s, f_{\phi}(\epsilon; s))$

```
73     # update actor
74     mu, std = actor(torch.Tensor(states))
75     policy, log_policy = eval_action(mu, std)
76
77     q_value1, q_value2 = critic(torch.Tensor(states), policy)
78     min_q_value = torch.min(q_value1, q_value2)
79
80     actor_loss = ((alpha * log_policy) - min_q_value).mean() # Equation 9
81     actor_optimizer.zero_grad()
82     actor_loss.backward()
83     actor_optimizer.step()
```

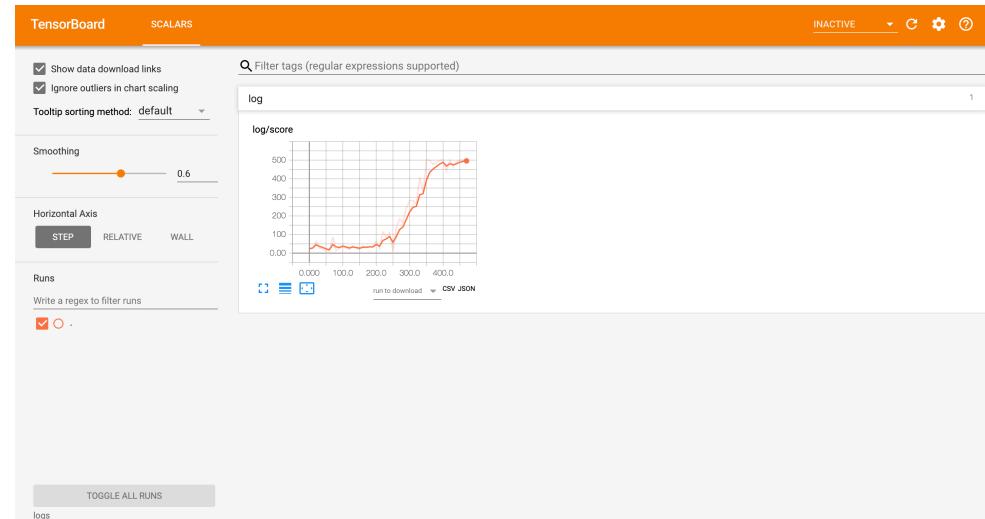
Train model

- Update alpha
 - `log_alpha` - (1)
 - `log_policy` - (64, 1)
 - `target_entropy` → -0.0
 - $J(\alpha) = -\frac{1}{N} \sum \alpha (\log \pi_\phi(f_\phi(\epsilon; s) | s) + \bar{\mathcal{H}})$
 - `alpha` - (1)

```
85     # update alpha
86     alpha_loss = -(log_alpha * (log_policy + target_entropy).detach()).mean() # Equation 18
87     alpha_optimizer.zero_grad()
88     alpha_loss.backward()
89     alpha_optimizer.step()
90
91     alpha = torch.exp(log_alpha)
92
93     return alpha
```

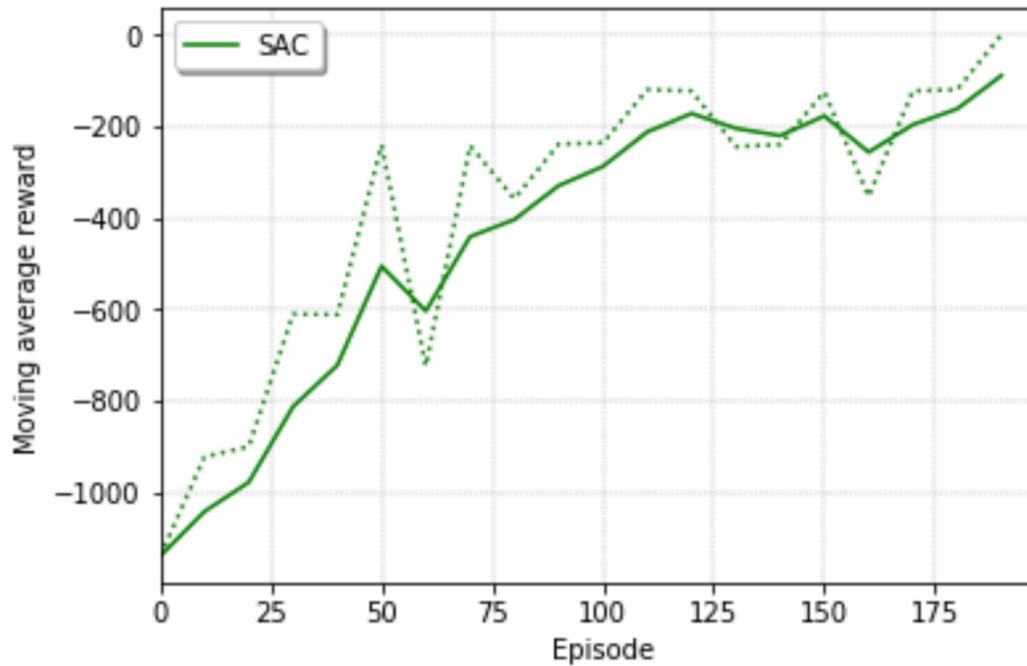
Train & TensorboardX

- Terminal A - train 실행
 - conda activate env_name
 - python train.py
- Terminal B - tensorboardX 실행
 - conda activate env_name
 - tensorboard --logdir logs
 - (웹에서) localhost:6006



Learning curve & Test

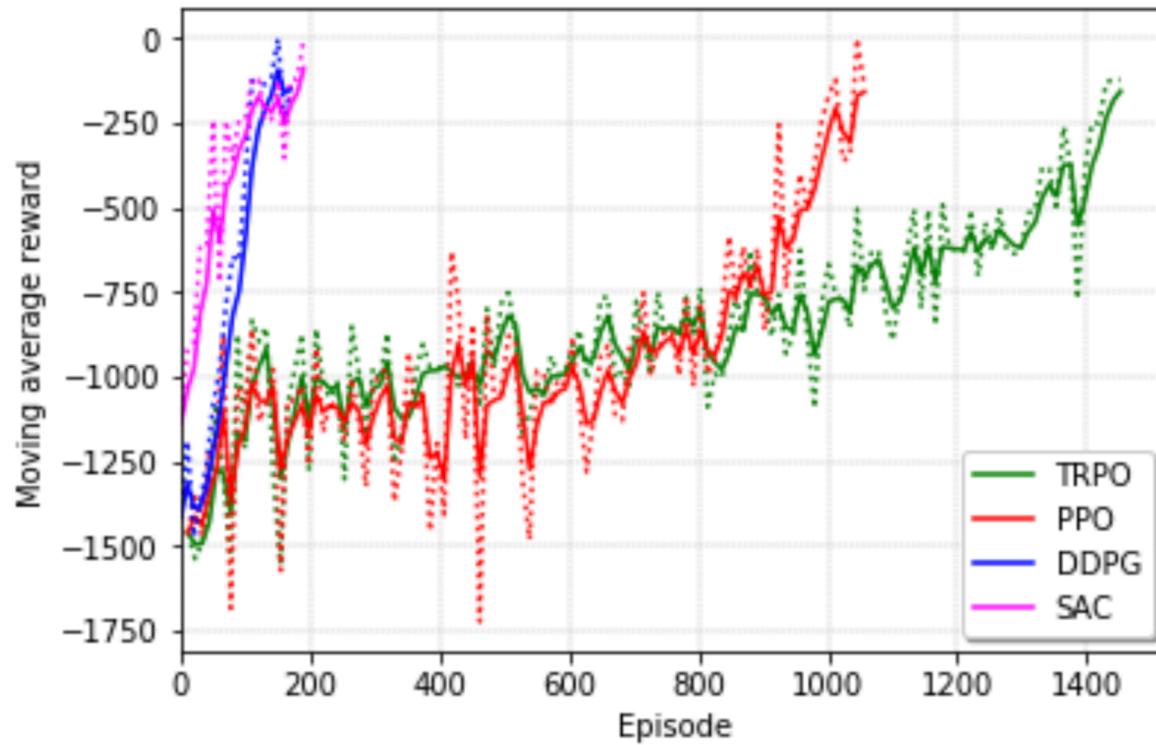
- Learning curve



- Test
 - `python test.py`

Comparison of algorithms for continuous action

- Learning curve



Comparison of algorithms for continuous action

Continuous action일 때는 어떤 알고리즘이 제일 좋을까요?
그리고 그 알고리즘이 **왜** 좋을까요?

- TRPO, PPO (On-policy algorithms)
 - Poor sample efficiency → 업데이트를 하기 위해서는 새로운 sample들을 수집해야 함
 - The number of gradient step and samples per step that are **extravagantly expensive**
- DDPG (Off-policy algorithm)
 - The interplay between the **deterministic actor** and the Q-function typically makes DDPG **extreme difficult to stabilize and brittle to hyperparameter settings**
 - Exploration noise can cause **sudden failures in unstable environments**
- SAC (Off-policy algorithm)
 - The **stochastic actor** aims to maximize expected reward while also **maximizing entropy**
 - Automatic gradient-based **temperature hyperparameter tuning** method

Thank you



CORE
Control + Optimization Research Lab